

U-HAP: Unified Authorization over Heterogeneous Policies with Efficient Multi-Resource Verification in Kubernetes

Krittapak Jairak*, Singha Junchan*, Phisitphon Pruksorranan*, Somchart Fugkeaw*

*School of Information, Computer and Communication Technology (ICT)

Sirindhorn International Institute of Technology, Thammasat University, Thailand

6622782207@g.siit.tu.ac.th, 6622770350@g.siit.tu.ac.th, 6622782322@g.siit.tu.ac.th, somchart@siit.tu.ac.th

Abstract—Kubernetes environments increasingly host multiple applications across pods and clusters, each governed by independently defined authorization policies such as RBAC, ABAC, and ACL. This heterogeneity subjects users to repeated and fragmented authorization checks when accessing multiple resources, resulting in increased latency, inconsistent decisions, and management complexity. While Single Sign-On (SSO) simplifies authentication, it does not address the challenge of unified authorization across heterogeneous policies. We propose U-HAP, a unified authorization framework for efficient and consistent multi-resource authorization across heterogeneous policies without requiring standardization. U-HAP introduces a universal graph-based representation that captures diverse policy models within unified semantics, reducing authorization to a reachability problem. By compiling resource-scoped policy artifacts and performing lightweight evaluation at request time, U-HAP eliminates redundant checks and ensures deterministic conflict resolution. We implement U-HAP as a Kubernetes webhook authorizer and evaluate it at scale. Experimental results demonstrate that U-HAP maintains near-constant authorization latency even as the number of namespaces scales, significantly outperforming traditional approaches. Furthermore, compilation-based optimization enables sub-linear scaling for complex policies, while decision caching provides additional speedups for repeated requests. U-HAP further demonstrates faster *policy-update* propagation than Open Policy Agent (OPA), completing the end-to-end edit-to-first-decision cycle significantly faster at scale under semantically equivalent policies verified by an automated decision-equivalence gate.

Index Terms—access control, Kubernetes, policy unification, directed acyclic graph, hash consing, RBAC, ABAC, ACL, webhook authorization, Zero Trust

I. INTRODUCTION

Kubernetes is the de facto platform for orchestrating containerized applications, including security-critical systems [1]. Yet authorization remains a major challenge: misconfigurations and excessive permissions frequently lead to vulnerabilities and full cluster compromise [2], [3].

Authorization in Kubernetes is inherently *multi-application*, *multi-resource*. A single workflow may touch multiple pods, namespaces, or clusters, each governed by independently defined policies (RBAC [4], ABAC, ACL [5], or custom) that operate under incompatible semantics. This results in repeated checks, inconsistent decisions, and operational inefficiency, even when SSO reduces authentication overhead.

Existing mechanisms suffer three key limitations: (i) *semantic fragmentation*—RBAC, ABAC, and ACLs use incompatible languages and semantics, increasing duplication and misconfiguration risk [6]; (ii) *implicit conflict resolution*—ordering across models yields non-deterministic decisions; and (iii) *runtime inefficiency*—repeated policy scanning evaluates irrelevant rules. Existing solutions remain insufficient for scalable *multi-model authorization* [7], [8].

We propose U-HAP, a *compilation-driven, index-based* framework that shifts computation offline: policies are compiled into resource–action scoped indices enabling a *two-level pruning* strategy—*policy-type pruning* skips inactive classes, while *token-driven pruning* limits evaluation to rules matching the caller’s roles and attributes. Although implemented as a Kubernetes webhook, the design is general and applicable to other distributed environments. This paper contributes: (i) a unified graph-based model integrating RBAC, ABAC, ACL, and deny; (ii) a compilation-driven architecture eliminating runtime policy scanning; (iii) a lightweight two-level pruning strategy; and (iv) a Kubernetes webhook implementation demonstrating low-latency, deterministic authorization with unified semantics.

II. RELATED WORK

Security and access control in Kubernetes environments have been widely studied, particularly from the perspectives of system hardening and configuration correctness. Foundational guidelines such as the NSA/CISA Kubernetes Hardening Guidance [1] and subsequent compliance-oriented analyses [9] emphasize secure control plane configuration, workload isolation, and least-privilege enforcement. However, empirical studies reveal that misconfigurations and excessive permissions remain pervasive in real-world deployments [2], [3], often enabling privilege escalation and even full cluster compromise. Recent tools such as EPScan [10] further confirm that excessive RBAC permissions in third-party applications continue to be a systematic security risk, while KubeKeeper [19] mitigates this risk via a sub-RBAC, per-pod admission webhook for Kubernetes secrets—the same extension point U-HAP uses, but restricted to a single resource type. Complementary work on formal verification [11] demonstrates that static policy analysis can prevent misconfigurations in Kubernetes clusters, yet such

approaches address correctness of individual policies rather than unified authorization across mixed-model environments.

Kubernetes natively adopts Role-Based Access Control (RBAC) [4], which provides scalable permission management but lacks expressiveness for fine-grained, context-aware authorization. KubeFence [20] addresses this by compiling per-workload API filters at deploy time—aligned with U-HAP’s compile-then-enforce philosophy but operating at the container-image level rather than as a unified ACL/RBAC/ABAC/Deny authorizer. To address these limitations, policy-as-code frameworks enable declarative and dynamic policy enforcement, while Zero Trust architectures [12] promote continuous verification and context-aware access decisions in distributed cloud-native environments. Nevertheless, these approaches operate on independently evaluated policy models and require repeated evaluation across multiple enforcement layers, leading to redundant checks and increased latency in multi-resource authorization scenarios [6]. Recent optimization efforts such as PerfSPEC [13] improve enforcement efficiency through performance-aware profiling, and compile-time policy optimization techniques [14] reduce redundant rule evaluation. However, both approaches operate within individual policy models and do not address unified authorization across diverse policy types.

Beyond RBAC, Attribute-Based Access Control (ABAC) and standardized frameworks such as XACML [7] provide more expressive and fine-grained authorization capabilities. However, XACML relies on heavyweight policy decision point (PDP) architectures that are not well suited for lightweight, low-latency cloud-native systems. Prior work on compiling XACML policies into indexed evaluation structures [15] has demonstrated significant speedups over sequential evaluation, motivating the compile-time optimization approach adopted by U-HAP. Recent works explore graph-based representations of access control, modeling authorization as relationship inference over entities and attributes [8], [16]. Large-scale systems such as Zanzibar [17] demonstrate the scalability of relationship-based authorization, but require runtime graph traversal, introducing overhead proportional to graph complexity. Similarly, approaches that convert ABAC policies into RBAC models [18] improve interoperability but do not resolve the challenge of unified and efficient runtime authorization across diverse policy types.

Overall, existing works address individual aspects of Kubernetes authorization—security hardening [1], [9], policy expressiveness [4], [7], compile-time optimization [14], [15], or scalable enforcement [17]—but none provides a unified framework that simultaneously integrates RBAC, ABAC, ACL, and explicit deny under a single compilation model with deterministic runtime evaluation. U-HAP bridges this gap by treating heterogeneous policy models as first-class citizens of a single compilation pipeline, yielding expressive coverage, predictable $O(1)$ per-request latency, and scalable policy-update throughput.




Fig. 1. U-HAP compilation-driven architecture on Kubernetes.

III. PROPOSED SYSTEM

A. System Model

U-HAP separates *policy compilation* (offline) from *request-time authorization* (online) to achieve low-latency, deterministic enforcement (Fig. 1).

System Entities.

- **User ($\mathcal{U}$):** Authenticated subject issuing requests with dynamic attributes (roles, groups, context).
- **Administrator:** Defines policies (RBAC, ABAC, ACL, Deny) via the U-HAP DSL.
- **Kubernetes API Server ($\mathcal{K}$):** Dispatches authorization queries via `SubjectAccessReview` (SAR).
- **U-HAP Webhook ($\mathcal{W}$):** Evaluates SAR requests using compiled artifacts and returns decisions.
- **DSL Loader ($\mathcal{L}$):** Parses and validates policies.
- **Policy Compiler ($\mathcal{B}$):** Compiles policies into resource–action scoped artifacts $\{\mathcal{C}_{n,r,a}\}$.
- **Graph Registry ($\mathcal{M}$):** In-memory store for compiled artifacts, enabling $O(1)$ lookup.
- **Resources ($\mathcal{O}$):** Protected Kubernetes objects organized by namespace.

U-HAP decouples authentication from authorization by deferring user-specific bindings to request time rather than embedding them in precomputed structures. Identity, roles, and attributes are supplied via tokens at request time, enabling dynamic evaluation without recompilation.

Trust Assumptions. The API server and U-HAP components are trusted. Policies are assumed correct after validation. Adversaries may issue malicious requests but cannot tamper with compiled artifacts or the registry.

B. Formal Graph Model

U-HAP defines a unified graph abstraction for heterogeneous authorization policies. For each namespace–resource pair (n, r) , the system constructs

$$G_{n,r} = (V_{n,r}, E_{n,r}) \quad (1)$$

where $V_{n,r}$ and $E_{n,r}$ capture all policy semantics for resource r .

The vertex set $V_{n,r}$ comprises typed nodes: subject nodes $s \in \mathcal{S}$ (authenticated identities), role nodes $\rho \in \mathcal{R}$ (RBAC abstractions), gate nodes $g \in \mathcal{G}$ (ABAC predicates), deny nodes $d \in \mathcal{D}$ (explicit denials), and a distinguished resource node r . The edge set $E_{n,r}$ encodes relationships: (s, ρ) for subject–role assignments, (ρ_i, ρ_j) for role hierarchy, (ρ, r) for role-based permissions, (s, r) for direct ACL permissions, (s, g) and (g, r) for attribute-based constraints, and (s, d) , (d, r) for explicit deny semantics.

Well-formedness requires that (i) the role hierarchy is acyclic, (ii) all edges terminate at a gate, deny, or resource node, and (iii) deny nodes have highest precedence, ensuring deterministic conflict resolution.

Authorization is defined as a constrained reachability problem over $G_{n,r}$. Given a request (s, a, r) with context $\mathbf{x}$, access is denied if a path $s \rightarrow d \rightarrow r$ exists; otherwise, access is granted if: (i) a direct edge (s, r) exists (ACL), (ii) a path $s \rightarrow \rho^* \rightarrow r$ exists through the role hierarchy (RBAC), or (iii) a path $s \rightarrow g \rightarrow r$ exists such that $g(\mathbf{x})$ is true (ABAC). If none apply, access is denied by default.

C. System Process

The proposed U-HAP system consists of three major phases as follows:

Phase 1: Setup (Step-Based Description): The setup phase initializes the U-HAP system and is executed once during deployment. It establishes the foundational structures required for policy compilation and runtime authorization.

Step 1: System Initialization.

Define the set of namespaces $\mathcal{N}$ and, for each $n \in \mathcal{N}$, the set of protected resources $\mathcal{O}_n$.

Configure the Kubernetes API server to delegate authorization decisions via `SubjectAccessReview` (SAR) requests to the U-HAP Authorization Webhook. Normalize each incoming request into a unified representation:

$$\mathcal{Q} = \langle \tau, (n, r, a) \rangle \quad (2)$$

where τ is the user token and (n, r, a) denotes namespace, resource, and action.

Step 2: Component Deployment. U-HAP deploys four core components: the *DSL Loader*, which parses policy definitions; the *Policy Compiler/Graph Builder*, which transforms policies into graph-based structures; the *Graph Registry*, which stores compiled artifacts for constant-time lookup; and the *Authorization Webhook*, which handles runtime authorization requests. The DSL Loader and Compiler operate offline during policy preparation, while the Registry and Webhook serve the runtime authorization path.

Step 3: State Initialization. The system initializes both a *persistent policy store* (e.g., Kubernetes ConfigMaps) to maintain policy definitions and an *in-memory registry* to enable efficient runtime access. For each (n, r, a) , a structured container of compiled artifacts is allocated as

$$C_{n,r,a} = (G_{n,r}, I_{n,r,a}^{acl}, I_{n,r,a}^{rbac}, I_{n,r,a}^{abac}, I_{n,r,a}^{attr}, I_{n,r,a}^{deny}, \Psi_{n,r,a}), \quad (3)$$

where each component corresponds to a graph structure, policy index, or auxiliary metadata as defined in Table I.

Step 4: Token Trust Configuration.

Define the authenticated user token structure as:

$$\tau = \langle uid, \mathcal{R}_u, \mathcal{G}_u, \mathcal{A}_u, t_{exp}, \sigma \rangle \quad (4)$$

Where $uid, \mathcal{R}_u, \mathcal{G}_u, \mathcal{A}_u, t_{exp}, \sigma$ denote the authenticated user identity, the role set, the group set, the user attributes, the token expiration, and the cryptographic signature, respectively. Assume token authenticity is verified prior to authorization (e.g., via Kubernetes authentication layer). Define the supported policy types:

$$\Pi_n = \{\text{RBAC}, \text{ABAC}, \text{ACL}, \text{Deny}\}$$

with enforced schema and validation constraints.

Phase 2: Policy Compilation: The compilation phase transforms DSL policies into *resource–action–scoped runtime artifacts* consumed by the verifier in Phase 3. Rather than relying on runtime graph traversal, U-HAP compiles policies into verification-oriented indices enabling constant-time lookup, token-driven pruning, and short-circuit evaluation.

Input and Output. For $\Pi_n = \{\text{RBAC}, \text{ABAC}, \text{ACL}, \text{Deny}\}$, the compiler produces:

$$C_n = \{C_{n,r,a} \mid r \in \mathcal{O}_n, a \in \mathbb{A}\} \quad (5)$$

where each unit $C_{n,r,a}$ aligns with the verification input $\langle \tau, (n, r, a) \rangle$, eliminating runtime filtering.

Compilation Procedure.

Step 1: Ingestion and IR Construction. Policies are parsed into:

$$\mathcal{IR} = (\mathcal{S}, \mathcal{R}, \mathcal{P}, \mathbb{A}, \mathcal{D}) \quad (6)$$

Where $\mathcal{S}, \mathcal{R}, \mathcal{P}, \mathbb{A}, \mathcal{D}$ denote the sets of subjects, roles, permissions, actions, and deny rules, respectively.

Step 2: Validation and Normalization. The DSL Loader enforces syntax, type safety, reference consistency, and acyclic role hierarchies. Role inheritance is reduced to transitive closure $\rho_i \Rightarrow^* \rho_j$.

Step 3: Resource–Action Decomposition. Policies are partitioned as $\Pi_n \rightarrow \{\Pi_{n,r,a}\}$, ensuring each artifact corresponds exactly to (n, r, a) .

Step 4: Artifact Generation. For each (n, r, a) , the artifacts shown in Table I are constructed.

TABLE I
COMPILED ARTIFACTS PER RESOURCE–ACTION SCOPE (n, r, a) .

Artifact	Definition	Notes
Graph	$G_{n,r}$	audit only
Deny	$I_{n,r,a}^{deny} = \{(\text{scope}(d), \phi_d)\}$	constant-time match
ACL	$I_{n,r,a}^{acl} = \mathcal{S}_{r,a}^{acl}$	membership check
RBAC	$I_{n,r,a}^{rbac} = \{\rho \mid \rho \Rightarrow^* (r, a)\}$	$\mathcal{R}_u \cap I_{n,r,a}^{rbac} \neq \emptyset$
ABAC	$I_{n,r,a}^{abac} = \{(g_\phi, \text{req}(\phi))\}$	$g_\phi : \mathcal{A}_u \rightarrow \{0, 1\}$
Attr	$I_{n,r,a}^{attr} : k \mapsto \{\phi \mid k \in \text{req}(\phi)\}$	selective evaluation
Summary	$\Psi_{n,r,a} = \langle \mathcal{S}_{r,a}^{acl}, \mathcal{R}_{r,a}^{rbac}, \Phi_{r,a}^{abac}, \mathcal{D}_{r,a} \rangle$	–

where $\text{scope}(d)$ is the set of principals to which deny rule d applies, ϕ_d is the condition predicate of deny rule d , and $\text{req}(\phi)$ is the set of attribute keys required by predicate ϕ . Each ABAC predicate ϕ is compiled into a DAG of GateNode objects via *hash consing*: structurally identical sub-expressions across policies share a single node, so any predicate appearing in k policies maps to exactly one GateNode in the shared DAG.

Step 5: Finalization and Registration. Redundant edges are removed, unreachable nodes are pruned, and consistency is verified. All $\mathcal{C}_{n,r,a}$ are then stored in the Graph Registry for $O(1)$ retrieval. ACL reduces to membership checking, RBAC to set intersection, ABAC to attribute-key-pruned predicate evaluation, and deny to prioritized filtering, eliminating global policy scans at runtime.

Phase 3: Request-Time Authorization: At runtime, U-HAP evaluates each request using a *policy-type-pruned, token-projected, index-optimized* pipeline over the compiled artifacts generated in Phase 2 and summarized in Table I. Given a request $\mathcal{Q} = \langle \tau, (n, r, a) \rangle$, the goal is to determine whether the user represented by token τ is authorized to perform action a on resource r in namespace n .

Step 1: Request Reception and Token Extraction. The Authorization Webhook receives $\mathcal{Q} = \langle \tau, (n, r, a) \rangle$, where the validated token is

$$\tau = \langle uid, \mathcal{R}_u, \mathcal{G}_u, \mathcal{A}_u, t_{\text{exp}}, \sigma \rangle. \quad (7)$$

Here, uid denotes the authenticated user identity, $\mathcal{R}_u$ the role set, $\mathcal{G}_u$ the group set, and $\mathcal{A}_u$ the user and contextual attributes.

Step 2: Constant-Time Artifact Retrieval. Using the tuple (n, r, a) as a lookup key, the Webhook retrieves the compiled container $\mathcal{C}_{n,r,a}$ from the registry in $O(1)$ time. This container includes the graph, policy indices, and metadata defined in Table I. It also provides the fast-path activation descriptor

$$\Xi_{n,r,a} = \langle \lambda^{\text{deny}}, \lambda^{\text{acl}}, \lambda^{\text{rbac}}, \lambda^{\text{abac}} \rangle, \quad (8)$$

where each bit $\lambda^* \in \{0, 1\}$ indicates whether the corresponding policy type has at least one compiled rule for the requested (n, r, a) . Any class with $\lambda = 0$ is skipped entirely.

Step 3: Candidate Pruning. To avoid scanning irrelevant rules, U-HAP prunes the search space in two stages. First, *policy-type pruning* uses $\Xi_{n,r,a}$ to eliminate inactive policy classes. Second, *token-driven pruning* derives only the candidates relevant to the current user token:

$$\Gamma_{\text{deny}} = \{d \mid \text{scope}(d) \cap (\{uid\} \cup \mathcal{G}_u \cup \mathcal{R}_u) \neq \emptyset\}, \quad (9)$$

$$\Gamma_{\text{acl}} = (\{uid\} \cup \mathcal{G}_u) \cap I_{n,r,a}^{\text{acl}}, \quad (10)$$

$$\Gamma_{\text{rbac}} = \mathcal{R}_u \cap I_{n,r,a}^{\text{rbac}}, \quad (11)$$

$$\Gamma_{\text{abac}} = \bigcup_{k \in \text{dom}(\mathcal{A}_u)} I_{n,r,a}^{\text{attr}}(k). \quad (12)$$

Thus, Phase 3 evaluates only a compact request-specific subset rather than the full policy set.

Step 4: Decision Cache Check. Before evaluating rules, U-HAP checks a short-lived cache keyed by

$$\kappa = H(uid, n, r, a, \mathbf{b}_u, \mathbf{g}_u, \text{sig}(\mathcal{A}_u)), \quad (13)$$

where $\mathbf{b}_u$ and $\mathbf{g}_u$ are compact bit-vector encodings of user role and group memberships, $H(\cdot)$ is a collision-resistant hash function mapping the tuple to a κ , and $\text{sig}(\mathcal{A}_u)$ is a deterministic digest over the user attribute map $\mathcal{A}_u$. Thus, if κ is found in the cache, the system immediately returns the cached decision (e.g., allow) without further evaluation.

Step 5: Ordered Policy Verification. On a cache miss, U-HAP evaluates policies in fixed order $\text{Deny} \rightarrow \text{ACL} \rightarrow \text{RBAC} \rightarrow \text{ABAC} \rightarrow \text{Default Deny}$, terminating upon a definitive decision.

(a) *Deny Check.* Deny rules take precedence. If any matches, the request is rejected:

$$\exists d \in \Gamma_{\text{deny}} : \text{match}_{\text{deny}}(d, \tau, n, r, a) = 1 \Rightarrow \text{DENY}, \quad (14)$$

where $\text{match}_{\text{deny}} = 1$ iff $uid \in \text{scope}(d)$ and $\phi_d(\tau, n, r, a) = 1$. This evaluation order enforces a *monotone safety property*: no subsequent ACL, RBAC, or ABAC allow rule can override an explicit deny, providing a formal guarantee against privilege escalation via policy interaction.

(b) *ACL Check.* If no deny applies, ACL authorization is checked via constant-time hash membership:

$$uid \in I_{n,r,a}^{\text{acl}} \vee \mathcal{G}_u \cap I_{n,r,a}^{\text{acl}} \neq \emptyset. \quad (15)$$

If true, the request is allowed.

(c) *RBAC Check.* RBAC is evaluated using precompiled bit-vectors for efficient role matching. Let $\mathbf{b}_u$ encode the roles in $\mathcal{R}_u$, and let $\mathbf{b}_{n,r,a}^{\text{rbac}}$ denote the authorized-role bit-vector for (n, r, a) after applying transitive closure over the role hierarchy. Authorization holds iff:

$$\mathbf{b}_u \wedge \mathbf{b}_{n,r,a}^{\text{rbac}} \neq \mathbf{0}. \quad (16)$$

This constant-time bitwise operation eliminates the need for runtime role-hierarchy traversal and ensures scalable authorization.

(d) *ABAC Check.* If RBAC fails, only relevant ABAC predicates are evaluated, sorted by cost:

$$\Gamma_{\text{abac}}^\uparrow = \text{sort}(\Gamma_{\text{abac}}, \text{cost}). \quad (17)$$

Authorization holds if any predicate succeeds:

$$\exists \phi \in \Gamma_{\text{abac}}^\uparrow : g_\phi(\mathcal{A}_u) = 1 \Rightarrow \text{ALLOW}. \quad (18)$$

Composite predicates are:

$$g_{\phi_1 \wedge \phi_2} = g_{\phi_1} \wedge g_{\phi_2}, \quad g_{\phi_1 \vee \phi_2} = g_{\phi_1} \vee g_{\phi_2}, \quad (19)$$

$$g_{\text{th}(k, \{\phi_i\})} = \mathbb{1} \left[\sum_i g_{\phi_i} \geq k \right]. \quad (20)$$

Step 6: Default and Cache Update.

$\text{Verify}(\tau, n, r, a)$ returns a cached decision on hit. Otherwise, it applies the above order: $\text{Deny} \rightarrow \text{Allow}$

(ACL/RBAC/ABAC) $\rightarrow$ Default Deny, and caches the result under κ .

Runtime Complexity. Artifact retrieval is $O(1)$ by hash-keyed registry. ACL membership is $O(1)$ via hash set; RBAC is $O(1)$ via bitwise AND; ABAC is $O(|\text{gates}|)$ with per-node memoization guaranteeing each gate is evaluated at most once per request; deny pruning is $O(|\Gamma_{\text{deny}}|)$, bounded by applicable rules. Combined with the decision cache, Phase 3 is worst-case $O(|\text{gates}|)$ and typically $O(1)$ on warm requests— independent of total policy-set size or namespace count.

IV. EXPERIMENTAL EVALUATION

We evaluate U-HAP (Python 3.11) on an AMD Ryzen 9 7945HX (16C/32T), 32 GB DDR5-4800, CachyOS Linux. We compare against a conventional *SSO-based baseline* in which access rights are verified by evaluating application-specific policies stored in the cluster. Each data point is the median of 1,000 single-threaded iterations after 50 warm-ups, measuring only in-memory policy lookup and verification (no network or auth overhead). Each *namespace* models an independent application or workload.

Experiment 1 (Policy Verification Efficiency). We compare U-HAP, U-HAP with caching, and a conventional SSO baseline as the number of namespaces $N \in \{10, 50, 100, 200, 300, 500, 750, 1000\}$ grows. Each namespace contains 46 rules (10 RBAC, 20 ABAC, 10 ACL, 1 deny, 5 hierarchy edges); ABAC predicates use 10 logical gates (AND/OR/ATLEAST) and 11 attributes from a 20-key pool, with $\sim 50\%$ atom sharing across rules.




Fig. 2. Per-request authorization latency vs. total number of application namespaces.

Fig. 2 shows near-constant latency for all configurations as N scales from 10 to 1,000: namespace growth does not affect the verification cost of any target. U-HAP achieves $\approx 1.7\times$ lower latency than the baseline via compiled evaluation (shared ABAC DAG, hash-set ACL, bit-vector RBAC), while caching provides an additional $\sim 14\times$ speedup on repeated requests by bypassing evaluation entirely. In contrast, the SSO baseline performs sequential policy scanning—including repeated role resolution and hierarchy traversal for RBAC—

whose cost scales linearly with namespace count, explaining its proportionally higher and less predictable latency.

Experiment 2 (Policy Size Impact). We vary $k \in [5, 320]$ rules per model (one model type per namespace). ABAC predicates use 10 gates with 50% atom sharing; the matching rule sits at the median position so the baseline scans $\sim k/2$ rules. Caching is disabled to expose raw evaluation cost.




Fig. 3. Per-model authorization latency (log scale) vs. number of rules ($k=5-320$).

Fig. 3 shows compilation gains increase with k . ABAC benefits most: hash-consed DAG evaluation yields up to $16.9\times$ speedup at $k=320$ (crossover $k \approx 20$; minor $0.7\times$ overhead at $k=5$). ACL and RBAC remain near-constant via hash-set and bit-vector evaluation, reaching $5.0\times$ and $5.6\times$ at $k=320$. Overall, compilation overhead at very small k is marginal and quickly amortized; performance gains become significant beyond the crossover point and grow monotonically with rule-set scale.

Experiment 3 (Policy Update Latency vs. OPA). We measure end-to-end *update latency* (policy edit to first decision under the new policy) against OPA [21], the CNCF graduated general-purpose engine. Each U-HAP policy set is mechanically translated to equivalent Rego; an automated gate confirms 100/100 identical decisions at every $n \in [10, 2000]$. Policy sets ($n \in \{10, 100, 500, 1000, 2000\}$, fixed mix: 40% RBAC, 30% ABAC, 15% ACL, 10% deny, 5% hierarchy) run 30 timed updates per n after 3 warm-ups; we report medians with 95% bootstrap CIs and additionally post-parse (engine-only) latency, since YAML parsing is a fixed cost both engines pay.

As shown in Table II and Fig. 4, U-HAP completes the update cycle faster at every n , from $1.63\times$ at $n=10$ to $2.73\times$ at $n=2000$. The post-parse advantage grows to $19.5\times$ at $n=2000$ (31 ms vs. 611 ms), demonstrating the benefit of compile-time hash consing over OPA's per-update Rego compilation. An OPA bundle-activation mode lands within $\pm 4\%$ of REST PUT, ruling out a deployment-mode artefact. The compiled-state footprint stays ≤ 2.7 MiB throughout.

V. CONCLUSION

This paper presented U-HAP, a compilation-driven authorization framework that unifies RBAC, ABAC, ACL, and




Fig. 4. End-to-end update latency vs. policy count (log-log). Bars are 95% percentile-bootstrap CIs.

TABLE II

UPDATE LATENCY MEDIANS (MS). POST-PARSE EXCLUDES YAML PARSE (ENGINE-AGNOSTIC COST). SPEEDUP = OPA / U-HAP.

n	Total (ms)			Post-parse (ms)		
	U-HAP	OPA	×	U-HAP	OPA	×
10	1.75	2.85	1.63×	0.14	2.85	20.4×
100	14.76	26.60	1.80×	2.03	26.60	13.1×
500	57.35	119.40	2.08×	7.87	119.40	15.2×
1000	109.99	277.32	2.52×	16.27	277.32	17.0×
2000	223.41	610.54	2.73×	31.31	610.54	19.5×

explicit deny policies into a single Kubernetes authorization mechanism. By shifting policy reasoning to an offline compilation phase, U-HAP eliminates repeated runtime policy scanning and enables efficient authorization through indexed artifacts and lightweight verification. Experimental results demonstrate that U-HAP maintains stable low latency under increasing namespace and policy complexity, consistently outperforming conventional SSO-based approaches that rely on repeated policy evaluation. Compilation and decision caching provide complementary performance benefits, while per-policy analysis confirms improved efficiency across RBAC, ABAC, and ACL models. U-HAP also achieves faster policy-update processing with a compact compiled-state footprint, demonstrating the practicality of the framework for dynamic Kubernetes environments. Future work includes supporting cluster-wide and cross-namespace policies, enabling incremental recompilation for dynamic updates, and integrating U-HAP with production-grade authorization engines for large-scale multi-cluster deployment and evaluation.

REFERENCES

- [1] NSA and CISA, “Kubernetes Hardening Guidance,” NSA Cybersec. Tech. Report, ver. 1.2, Aug. 2022. [Online]. Available: https://media.defense.gov/2022/Aug/29/2003066362/-1/-1/0/CTR_KUBERNETES_HARDENING_GUIDANCE_1.2_20220829.PDF
- [2] A. Rahman, S. I. Shamim, D. B. Bose, and R. Pandita, “Security Misconfigurations in Open Source Kubernetes Manifests: An Empirical Study,” *ACM Trans. Softw. Eng. Methodol.*, vol. 32, no. 4, 2023, doi: 10.1145/3579639.

- [3] N. Yang, W. Shen, J. Li, X. Liu, X. Guo, and J. Ma, “Take Over the Whole Cluster: Attacking Kubernetes via Excessive Permissions of Third-Party Applications,” in *Proc. ACM CCS*, 2023, pp. 3048–3062, doi: 10.1145/3576915.3623121.
- [4] G. Rostami, “Role-Based Access Control (RBAC) Authorization in Kubernetes,” *J. ICT Standardization*, vol. 11, no. 3, pp. 237–260, 2023, doi: 10.13052/jicts2245-800X.1132.
- [5] N. Farhadighalati, L. A. Estrada-Jimenez, S. Nikghadam-Hojjati, and J. Barata, “A Systematic Review of Access Control Models: Background, Existing Research, and Challenges,” *IEEE Access*, vol. 13, pp. 17777–17806, 2025, doi: 10.1109/ACCESS.2025.3533145.
- [6] M. S. Rahaman, S. N. Tisha, E. Song, and T. Cerny, “Access Control Design Practice and Solutions in Cloud-Native Architecture: A Systematic Mapping Study,” *Sensors*, vol. 23, no. 7, art. no. 3413, 2023, doi: 10.3390/s23073413.
- [7] OASIS, “eXtensible Access Control Markup Language (XACML) Version 3.0 Plus Errata 01,” OASIS Standard, 2017.
- [8] M. Yang, V. Atluri, S. Sural, and J. Vaidya, “A Graph-Based Framework for ABAC Policy Enforcement and Analysis,” in *Proc. DBSec*, 2024, pp. 3–23, doi: 10.1007/978-3-031-65172-4_1.
- [9] Z. Moric, V. Dakić, and T. Čavala, “Security Hardening and Compliance Assessment of Kubernetes Control Plane and Workloads,” *J. Cybersec. Privacy*, vol. 5, no. 2, art. no. 30, 2025, doi: 10.3390/jcp5020030.
- [10] Y. Gu, X. Tan, Y. Zhang, S. Gao, and M. Yang, “EPScan: Automated Detection of Excessive RBAC Permissions in Kubernetes Applications,” in *Proc. 2025 IEEE Symp. Security and Privacy (SP)*, 2025, pp. 3199–3217, doi: 10.1109/SP61157.2025.00011.
- [11] A. Sissodiyi, E. Chiquito, U. Bodin, and J. Kristiansson, “Formal Verification for Preventing Misconfigured Access Policies in Kubernetes Clusters,” *IEEE Access*, vol. 13, pp. 141798–141813, 2025, doi: 10.1109/ACCESS.2025.3597504.
- [12] R. Chandramouli and Z. Butcher, “A Zero Trust Architecture Model for Access Control in Cloud-Native Applications in Multi-Location Environments,” NIST SP 800-207A, 2023, doi: 10.6028/NIST.SP.800-207A.
- [13] H. Kermabon-Bobinnec *et al.*, “PerfSPEC: Performance Profiling-Based Proactive Security Policy Enforcement for Containers,” *IEEE Trans. Dependable Secure Comput.*, vol. 22, no. 2, pp. 919–938, 2025, doi: 10.1109/TDSC.2024.3420712.
- [14] S. Kern, T. Baumer, S. Groll, L. Fuchs, and G. Pernul, “Optimization of Access Control Policies,” *J. Inf. Secur. Appl.*, vol. 70, art. no. 103301, 2022, doi: 10.1016/j.jisa.2022.103301.
- [15] A. X. Liu, F. Chen, J. H. Hwang, and T. Xie, “Designing Fast and Scalable XACML Policy Evaluation Engines,” *IEEE Trans. Computers*, vol. 60, no. 12, pp. 1802–1817, Dec. 2011, doi: 10.1109/TC.2010.274.
- [16] L. Ma *et al.*, “Research on Authorization Model of Attribute Access Control Based on Knowledge Graph,” in *Proc. UbiSec*, Springer, 2024, pp. 350–364, doi: 10.1007/978-981-97-1274-8_23.
- [17] R. Pang *et al.*, “Zanzibar: Google’s Consistent, Global Authorization System,” in *Proc. USENIX ATC*, 2019, pp. 33–46. [Online]. Available: <https://www.usenix.org/conference/atc19/presentation/pang>
- [18] M. Davari and M. Zulkernine, “Automatic Conversion of ABAC Policies for RBAC Systems,” in *Proc. 2023 IEEE Conf. Dependable and Secure Computing (DSC)*, 2023, pp. 1–7, doi: 10.1109/DSC61021.2023.10354157.
- [19] M. Rostamipoor, A. Sadeghi, and M. Polychronakis, “KubeKeeper: Protecting Kubernetes Secrets Against Excessive Permissions,” in *Proc. 2025 IEEE Eur. Symp. Security and Privacy (EuroS&P)*, 2025, pp. 322–338, doi: 10.1109/EuroSP63326.2025.00027.
- [20] C. Cesarano and R. Natella, “KubeFence: Security Hardening of the Kubernetes Attack Surface,” in *Proc. 2025 IEEE/IFIP Int. Conf. on Dependable Systems and Networks (DSN)*, 2025, pp. 497–510, doi: 10.1109/DSN64029.2025.00054.
- [21] Open Policy Agent, “OPA — Open Policy Agent,” CNCF graduated project, 2024. [Online]. Available: <https://www.openpolicyagent.org/>